



# PROJECT #2 - B+TREE

## OVERVIEW

The second programming project is to implement an **index** in your database system. The index is responsible for fast data retrieval without having to search through every row in a database table, providing the basis for both rapid random lookups and efficient access of ordered records.

You will need to implement [B+Tree](#) dynamic index structure. It is a balanced tree in which the internal pages direct the search and leaf pages contains actual data entries. Since the tree structure grows and shrink dynamically, you are required to handle the logic of split and merge. The project is comprised of the following three tasks:

- [Task #1 - B+Tree Pages](#)
- [Task #2 - B+Tree Data Structure](#)
- [Task #3 - Index Iterator](#)

This is a single-person project that will be completed individually (i.e., no groups).

**?** Please post all of your questions about this project on Canvas. Do **not** email the TAs directly with questions. The instructor and TAs will **not** teach you how to debug your code.

**Release Date:** Oct 02, 2017

**Due Date:** Oct 25, 2017 @ 11:59pm

## PROJECT SPECIFICATION

Like the first project, we are providing you with stub classes that contain the API that you need to implement. You should **not** modify the signatures for the pre-defined functions in these classes. If you do this, then it will break the test code that we will use to grade your assignment you end up getting no credit for the project. If a class already contains certain member variables, you should **not** remove them. But you may add private helper functions/member variables to these classes in order to correctly realize the functionality.

The correctness of B+Tree index depends on the correctness of your implementation of buffer pool, we will **not** provide solutions for the previous programming projects. Since the second project is closely related to the third project in which you will implement index crabbing within existing B+ index, we have passed in a pointer parameter called **transaction** with default value **nullptr**. You can safely ignore the parameter for now; you don't need to change or call any functions relate to the parameter.

- [B+Tree Parent Page](#)
- [B+Tree Internal Page](#)
- [B+Tree Leaf Page](#)

## B+TREE PARENT PAGE

This is the parent class that both the Internal Page and Leaf Page inherited from and it only contains information that both child classes share. The Parent Page is divided into several fields as shown by the table below. All multi-byte fields are stored with the most significant byte first (little-endian).

### B+Tree Parent Page Content

Variable Name	Size	Description
type	4	Page Type (internal or leaf)
size	4	Number of Key & Value pairs in page
max_size	4	Max number of Key & Value pairs in page
parent_id	4	Parent Page Id
page_id	4	Self Page Id

You must implement your Parent Page in the designated files. You are only allowed to modify the header file (`src/include/page/b_plus_tree_page.h`) and its corresponding source file (`src/page/b_plus_tree_page.cpp`).

## B+TREE INTERNAL PAGE

Internal Page does not store any real data, but instead it stores an ordered **m** key entries and **m+1** child pointers (a.k.a `page_id`). Since the number of pointer does not equal to the number of key, the first key is set to be invalid, and lookup methods should always start with the second key. At any time, each internal page is at least half full. During deletion, two half-full pages can be joined to make a legal one or can be redistributed to avoid merging, while during insertion one full page can be split into two.

You must implement your Internal Page in the designated files. You are only allowed to modify the header file (`src/include/page/b_plus_tree_internal_page.h`) and its corresponding source file (`src/page/b_plus_tree_internal_page.cpp`).

## B+TREE LEAF PAGE

The Leaf Page stores an ordered **m** key entries and **m** value entries. In your implementation, value should only be 64-bit `record_id` that is used to locate where actual tuples are stored, see **RID** class defined under in `src/include/common/rid.h`. Leaf pages have the same restriction on the number of key/value pairs as Internal pages, and should follow the same operations of merge, redistribute and split.

You must implement your Internal Page in the designated files. You are only allowed to modify the header file (`src/include/page/b_plus_tree_leaf_page.h`) and its corresponding source file (`src/page/b_plus_tree_leaf_page.cpp`).

`max_size` within both Internal Page's and Leaf Page's `Init()` method.

Each B+Tree leaf/internal page corresponds to the content (i.e., the byte array `data_`) of a memory page fetched by buffer pool. So every time you try to read or write a leaf/internal page, you need to first **fetch** the page from buffer pool using its unique `page_id`, then [reinterpret cast](#) to either a leaf or an internal page, and unpin the page after any writing or reading operations.

## TASK #2 - B+TREE DATA STRUCTURE

Your B+Tree Index could only support **unique key**. That is to say, when you try to insert a key-value pair with duplicate key into the index, it should not perform the insertion and return false.

Your B+Tree Index is required to correctly perform split if insertion triggers current number of key/value pairs exceeds `max_size`. It also needs to correctly perform merge or redistribute if deletion cause certain page to go below the occupancy threshold. Since any write operation could lead to the change of `root_page_id` in B+Tree index, it is your responsibility to update `root_page_id` in the header page (`src/include/page/header_page.h`) to ensure that the index is durable on disk. Within the `BPlusTree` class, we have already implemented a function called `UpdateRootPageId()` for you; all you need to do is invoke this function whenever the `root_page_id` of B+Tree index changes.

Your B+Tree implementation must hide the details of the key/value type and associated comparator, like this:

```
template <typename KeyType,
          typename ValueType,
          typename KeyComparator>
class BPlusTree{
    // ---
};
```

These classes are already implemented for you:

- **KeyType**: The type of each key in the index. This will only be `GenericKey`, the actual size of `GenericKey` is specified and instantiated with a template argument and depends on the data type of indexed attribute.
- **ValueType**: The type of each value in the index. This will only be 64-bit RID.
- **KeyComparator**: The class used to compare whether two `KeyType` instances are less/greater-than each other. These will be included in the `KeyType` implementation files.

## TASK #3 - INDEX ITERATOR

You will build a general purpose index iterator to retrieve all the leaf pages efficiently. The basic idea is to organize them into a single linked list, and then traverse every key/value pairs in specific direction stored within the B+Tree leaf pages. Your index iterator should follow the functionality of [Iterator](#) defined in c++11, including the ability to iterate through a range of elements using a set of operators (with at least the increment and dereference operators).

You must implement your index iterator in the designated files. You are only allowed to modify the header file (`src/include/index/index_iterator.h`) and its corresponding source file (`src/index/index_iterator.cpp`). You do not need to modify any other files. You **must** implement the

- `isEnd()` : Return whether this iterator is pointing at the last key/value pair.
- `operator++()` : Move to the next key/value pair.
- `operator*()` : Return the key/value pair this iterator is currently pointing at.

## INSTRUCTIONS

### SETTING UP YOUR DEVELOPMENT ENVIRONMENT

Download the project source code [here](#). This version has additional files that were added after project #1 so you need to update your local copy.

Make sure that your source code has the following `VERSION.txt` file:

```
Created: Oct 02 2017 @ 22:32:09
Last Commit: 250b733e4b16493caea9f3310dfefbf1029fdceae
```

### TESTING

You can test the individual components of this assignment using our testing framework. We use [GTest](#) for unit test cases. You can compile and run each test individually from the command-line:

```
cd build
make b_plus_tree_test
./test/b_plus_tree_test
```

In the `b_plus_tree_print_test`, you can print out the internal data structure of your b+ tree index, it's an intuitive way to track and find early mistakes.

```
cd build
make b_plus_tree_print_test
./test/b_plus_tree_print_test
```

**Important:** These tests are only a subset of the all the tests that we will use to evaluate and grade your project. You should write additional test cases on your own to check the complete functionality of your implementation.

### DEVELOPMENT HINTS

Instead of using `printf` statements for debugging, use the `LOG_*` macros for logging information like this:

```
LOG_INFO("# Leaf Page: %s", leaf_page->ToString().c_str());
LOG_DEBUG("Fetching page %d", page_id);
```

To enable logging in your project, you will need to reconfigure it like this:

```
cd build
cmake -DCMAKE_BUILD_TYPE=DEBUG ..
make
```

`LOG_LEVEL_INFO` (e.g., `LOG_INFO`, `LOG_WARN`, `LOG_ERROR`) will emit logging information.

Using `assert` to force check the correctness of your index implementation. For example, when you try to delete a page, the `page_count` must equals to 0.

Using a relatively small value of page size at the beginning test stage, it would be easier for you to check whether you have done the delete and insert in the correct way. You can change the page size in configuration file (`src/include/common/config.h`).

## GRADING RUBRIC

---

Each project submission will be graded based on the following criteria:

Does the submission successfully execute all of the test cases and produce the correct answer?

Does the submission execute without any memory leaks?

Note that we will use additional test cases that are more complex and go beyond the sample test cases that we provide you.

## LATE POLICY

---

See the [late policy](#) in the syllabus.

## SUBMISSION

---

After completing the assignment, you can submit your implementation of to Autolab:

<https://autolab.andrew.cmu.edu/courses/15445-f17>.

You only need to include the following files:

- Every file for Project 1 (6 in total)
- `src/include/page/b_plus_tree_page.h`
- `src/page/b_plus_tree_page.cp`
- `src/include/page/b_plus_tree_internal_page.h`
- `src/page/b_plus_tree_internal_page.cp`
- `src/include/page/b_plus_tree_leaf_page.h`
- `src/page/b_plus_tree_leaf_page.cp`
- `src/include/index/b_plus_tree.h`
- `src/index/b_plus_tree.cpp`
- `src/include/index/index_iterator.h`
- `src/index/index_iterator.cpp`

You can submit your answers as many times as you like and get immediate feedback. Your score will be sent via email to your Andrew account within a few minutes after your submission.

## COLLABORATION POLICY

---

Students are **not** allowed to copy the contents of a white-board after a group meeting with other students.  
Students are **not** allowed to copy the solutions from another colleague.

**⚠ WARNING:** All of the code for this project must be your own. You may not copy source code from other students or other sources that you find on the web. Plagiarism **will not** be tolerated. See CMU's [Policy on Academic Integrity](#) for additional information.

**Last Updated:** Mar 06, 2018

